




Kubernetes



LE TOUR DE TABLE

Pré-requis

Pour assister à cette formation, il est nécessaire de connaître et comprendre les notions de base associées aux conteneurs. Vous êtes capable de construire une image (par exemple avec un Dockerfile), lancer un conteneur, l'arrêter, inspecter ses logs.

- Un navigateur
- Un compte GCP (recommandé)



Agenda

- Rappel sur Docker
- Kubernetes : les origines
- Premier aperçu
- Architecture interne
- Démarrage du cluster
- Getting started
- Mise à l'échelle et mise à jour



Test de positionnement

<https://forms.gle/Q9h6iag4Q9snwJPr7>



Kubernetes : les origines



Google : 22 ans de containers

Borg

Gmail
Google Docs
Web Search

FlumeJava

Pregel



GFS/CFS

Bigtable

Megastore

MapReduce

Cloud Native Computing Foundation



Kubernetes, pour quoi faire ?

-
- Lancer 5 containers basés sur l'image redis: 4.0
 - Mettre en place un load-balancer interne au cluster pour servir ces 5 containers
 - Lancer 10 containers webapp: 1.0
 - Mettre en place un load-balancer public pour permettre d'accéder aux containers de l'extérieur du cluster
 - Augmenter le nombre de containers webapp pendant les soldes 😊
 - Continuer à servir les requêtes pendant la mise à jour vers webapp : 2.0




Mais aussi

- Mise à l'échelle automatique
- Déploiement “Blue/green” et “Canary”
- Exécution de traitements unitaires ou répétés
- Prioriser les tâches en cas de manque de ressources sur le cluster
- Exécuter des services qui persistent des données sur disque
- Contrôler l'accès aux différentes ressources
- Automatiser les tâches complexes (“operators”)

Premier aperçu




Cluster




[Ξ] School

Kubernetes cluster


Exécuter des applications



Node




Pod




Patterns


Sidecar




Adapter



Ambassador




Orchestrer les pods






Architecture interne




Master et Nodes




ApiServer et etcd



Kubelet




Controllers




Récapitulatif

- Le Kubelet s'enregistre dans la config
- L'utilisateur envoi l'état voulu sur etcd via l'apiserver
- Les contrôleurs et le Kubelet réagissent aux changements d'états pour obtenir et maintenir l'état voulu



À venir...



Getting started



Création de cluster Configuration de kubectl



Créer un cluster

GKE : Google Kubernetes Engine

[Minikube](#)



Utilitaires autour de kubectl

- <https://kubernetes.io/docs/tasks/tools/install-kubectl/#optional-kube-ctl-configuration>

```
$ kubectl <TAB> # autocompletion
```

- <https://github.com/ahmetb/kubectx>

```
$ kubectx gke-dev      # changer de contexte  
$ kubens kube-system  # modifier les namespaces du contexte courant
```

Utilitaires autour de kubectl

- ~800 kubectl aliases (bash/zsh)
<https://github.com/ahmetb/kubectl-aliases>
\$ kubectl # kubectl get pod
- Shell prompt
<https://github.com/jonmosco/kube-ps1>
(`* | minikube:default`) \$

Récupérer les exercices

- Clonez le dépôt github suivant:
git clone <https://github.com/raab-d/ESGI-K8S-5IABD1-2026>
- et placez-vous dans le dépôt cloné:
cd ESKI-K8S-5IABD1-2026/terraform/gke
- Démarrer un cluster K8S et configurer kubectl



Namespace

[Ξ] School

Namespace

- Espace de nom pour isoler les déploiements
- Peut être utilisé pour séparer les environnements (soft multi-tenant)
- Le nom d'une ressource est unique au sein d'un namespace.
- Par défaut:
 - default
 - kube-system
 - kube-public



Manifest

[Ξ] School

Manifest

```
kind: Pod
apiVersion: v1
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:alpine
status:
  ...
[☰] School
```

Créer et intéragir avec un pod



kubectl get pods

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-5dbbf858d-qfb7f	1/1	Running	0	2m

```
$ kubectl get pod -o yaml nginx-5dbbf858d-qfb7f
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-5dbbf858d-qfb7f
...

```

Pod yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:alpine
status:
  ...
[☰] School
```

Pods explained

\$ kubectl explain pod

KIND: Pod
VERSION: v1

DESCRIPTION:

Pod is a collection of containers that can run on a host. This resource is created by clients and scheduled onto hosts.

FIELDS:

```
apiVersion <string>
kind <string>
metadata <Object>
spec <Object>
[...]
```

[☰] School

Pods explained... deeper

\$ **kubectl explain pod.spec.containers.image**

KIND: Pod

VERSION: v1

FIELD: image <string>

DESCRIPTION:

Docker image name. More info:
<https://kubernetes.io/docs/concepts/containers/images> This field is optional to allow higher level config management to default or override container images in workload controllers like Deployments and StatefulSets.

Pod status (live)

...

status:

containerStatuses:

- **name: nginx**

containerID: docker://8db3af41eb87[...]3d103255

image: nginx:alpine

imageID: docker-pullable://nginx@sha256:1aed1[...]a0a5440

state:

running:

startedAt: 2018-08-13T21:08:10Z

hostIP: 192.168.65.3

podIP: 10.1.0.227

TP : Pod



Configuration d'une application



Configuration

- Configurer l'URI d'une base de données
- Injecter un fichier application-prod.yml
- Spécifier les credentials d'une api

Variables d'environnement

```
apiVersion: v1
kind: Pod
metadata:
  name: envar-demo
spec:
  containers:
    - name: envar-demo-container
      image: debian
  env:
    - name: DEMO_GREETING
      value: "Hello from the environment"
    - name: DEMO_FAREWELL
      value: "Such a sweet sorrow"
```

ConfigMap : création

```
$ kubectl create configmap <map-name> <data-source>  
$ kubectl create configmap my-app-conf  
--from-file=my-app-conf/configmap/application-dev.properties  
$ kubectl create configmap my-app-conf  
--from-literal=db-server=mydbserver.mycompany.com
```



ConfigMap: manifests

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-app-config
data:
  application.properties: |
    db-server=mydbserver.dev.mycompany.com
    username=my-rw-dbuser
$ kubectl apply -f configmap.yml
```

ConfigMap : utilisation

```
apiVersion: v1
kind: Pod
metadata:
  name: configmap-pod
spec:
  containers:
    - name: test
      image: busybox
    volumeMounts:
      - name: config-vol
        mountPath: /etc/config
          volumes:
            - name: config-vol
              configMap:
                name: log-config
                items:
                  - key: log_level
                    path: log_level
```

Point de montage du fichier: /etc/config/log_level

Secret : type

```
$ kubectl create secret --help
```

Available Commands:

- docker-registry Create a secret for use with a Docker registry
- generic Create a secret from a local file, directory or literal value
- tls Create a TLS secret



Secret : manifest

```
apiVersion: v1
kind: Secret
metadata:
name: mysecret
type: Opaque
data:
username: YWRtaW4=
password: MWYZDF1MmU2N2Rm
```

Secret : création

- Crée un secret depuis un fichier

```
$ kubectl create secret generic db-user-pass --from-file=./password.txt
```

- Crée un secret depuis une clé/valeur

```
$ kubectl create secret generic prod-db-secret  
--from-literal=username=produser
```

- Crée un secret "docker-registry "

```
$ kubectl create secret docker-registry regcred  
--docker-server=<your-registry-server> --docker-username=<your-name>  
--docker-password=<your-pword> --docker-email=<your-email>
```

Secret : env vars

```
apiVersion: v1
kind: Pod
spec:
  containers:
    - name: envar-demo-container
      image: debian
    env:
      - name: DEMO_GREETING
        valueFrom:
          secretKeyRef:
            name: demo      # le nom du secret
            key: password  # la clé dans le secret
```

Secret : volume

```
...  
kind: Pod  
spec:  
  containers:  
    - name: mypod  
      image: redis  
  volumes:  
    - name: secret-volume  
      secretName: mysecret  
  volumeMounts:  
    - name: secret-volume  
      mountPath: "/etc/foo"  
      readOnly: true
```


TP : Configuration




Service




Expose des pods




Service type ClusterIP




Service type NodePort




Service type LoadBalancer



Service type ExternalName



Service avec Endpoint explicite



Service yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  selector:
    app: nginx
  type: NodePort
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 80
```

Créer un service

```
$ kubectl apply -f service.yaml
```

- La ligne de commande suivante donne un résultat équivalent au fichier yaml :

```
$ kubectl expose nginx --port=80 --type=NodePort
```



Services

```
$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d
nginx	NodePort	10.97.47.155	<none>	80 31450 /TCP	2s

```
$ kubectl get svc -o yaml nginx
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
...

```

Service : DNS interne

- Entrée DNS : <service>.<namespace>.svc.cluster.local
- Et dans le namespace : <service>

TP : Service



ReplicaSet



ReplicaSet.yaml

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    <... pod definition ...>
```

kubectl get replicaset

```
$ kubectl get Replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-5dbbf858d	1	1	1	14m

```
$ kubectl get rs -o yaml nginx-5dbbf858d
```

```
apiVersion: extensions/v1beta1
kind: Replicaset
```

```
...
```

Replicaset status (live)

status:
⋮

```
availableReplicas: 1
fullyLabeledReplicas: 1
readyReplicas: 1
replicas: 1
```



Deployment



Deployment.yaml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:alpine
...
...
```

kubectl get deployment

```
$ kubectl get deployment  
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE  
nginx         1          1          1           1          1h
```

```
$ kubectl get deploy -o yaml nginx
```

```
apiVersion: extensions/v1beta1  
kind: Deployment
```

```
...
```

Mise à l'échelle manuelle

```
$ kubectl scale deployment/nginx --replicas=10
```

```
$ kubectl get all
```



The Label Game

[Ξ] School

Les labels Kubernetes

- Pour le fonctionnement interne de Kubernetes
 - ReplicaSet && Deployments ⇒ Pods
 - Services ⇒ Pods

“The label game”

App : MyApp
Phase : prod
Rôle : Interface



App : MyApp
Phase : prod
Rôle : BE



App : MyApp
Phase : test
Rôle : Interface



App : MyApp
Phase : test
Rôle : BE



[Ξ] school



“The label game”

App : MyApp
Phase : prod
Rôle : Interface



App : MyApp
Phase : test
Rôle : Interface



App : MyApp
Phase : prod
Rôle : BE




App : MyApp
Phase : test
Rôle : BE



“The label game”

App : MyApp
Phase : prod
Rôle : Interface



App : MyApp
Phase : test
Rôle : Interface




App : MyApp
Phase : prod
Rôle : BE



“The label game”

App : MyApp
Phase : prod
Rôle : Interface



App : MyApp
Phase : prod
Rôle : BE



App : MyApp
Phase : test
Rôle : Interface



App : MyApp
Phase : test
Rôle : BE



“The label game”

App : MyApp
Phase : prod
Rôle : Interface



App : MyApp
Phase : prod
Rôle : BE



App : MyApp
Phase : test
Rôle : Interface



App : MyApp
Phase : test
Rôle : BE




[Ξ] school



TP : Deployment









Rolling upgrade


Rolling upgrade 1/4




Rolling upgrade 2/4



Rolling upgrade 3/4



Rolling upgrade 4/4



Mise à jour **progressive**

- Mise à jour progressive (rolling-upgrade) sans interruption de service
- Gérée par le Deployment

Remplacer nginx par Apache httpd :

```
$ kubectl edit deployment nginx --record  
"image: nginx:alpine" ⇒ "image: httpd:alpine"
```



Annuler une mise à jour progressive

```
$ kubectl rollout history deployment nginx
```

REV CHANGE-CAUSE


- 1 kubectl apply --filename=deployment.yaml
- 2 kubectl edit deployment nginx
- 3 kubectl set image deploy nginx nginx=nginx:stable-alpine

```
$ kubectl rollout undo deployment --to-revision=1
```

Mise à l'échelle automatique



Horizontal Pod Autoscaler



Horizontal Pod Autoscaler

- Déclarer un Horizontal Pod scaler en CLI :

```
$ kubectl autoscale deployment nginx \
--cpu-percent=50 --min=1 --max=10
```



Horizontal Pod Autoscaler

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: helloworld
spec:
  maxReplicas: 10
  minReplicas: 1
  targetCPUUtilizationPercentage: 50
  scaleTargetRef:
    apiVersion: extensions/v1beta1
    kind: Deployment
    name: helloworld
```



[☰] School

[Ξ] School


Ingress



Ingress

- Point d'entrée unique du cluster
- HTTP (port 80) et HTTPS (port 443)
- Gestion des certificats SSL

Ingress



Ingress yaml

```
apiVersion: v1
kind: Ingress
metadata:
  name: nginx
spec:
  rules:
    - http:
        paths:
          - path: /
            backend:
              serviceName: nginx
              servicePort: http
```

Ingress

```
$ kubectl get ingress
```

NAME	HOSTS	ADDRESS	PORTS	AGE
nginx	*	localhost	80	3m

```
$ kubectl get ingress -o yaml nginx
```


```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: nginx
...

```

TP : Autoscaling

[Ξ] School

Récapitulatif



Stocker des données



Volumes

- Juste un dossier monté dans un/des containers
- Associé à la vie du Pod, survit au restart des containers
- Nombreuses implémentations :
 - emptyDir
 - gcePersistentDisk
 - configMap / secret
 - hostPath

emptyDir

- Volume vide, créé au démarrage d'un pod, supprimé avec la suppression du Pod.

volumes:

- name: cache-volume
- ```
emptyDir: {}
```

## **gcePersistentDisk**

- PersistentDisk GCE
  - Il doit être dans le même projet et la même zone que les VM du cluster
  - Il n'est pas supprimé à la suppression du pod



# **gcePersistentDisk**

volumes:

- name: test-volume

## **gcePersistentDisk:**

pdName: my-data-disk

fsType: ext4

# hostPath

- Monte dans le container un dossier du noeud sur lequel le Pod s'exécute

volumes:

- name: test-volume  
hostPath:  
# directory location on host  
path: /data  
# this field is optional  
type: Directory

# **persistentVolume & persistentVolumeClaim**

- L'admin crée une ressource PersistentVolume associée à une espace de stockage
- Le pod réclame du disque avec un PersistentVolumeClaim



# **persistentVolume**

```
kind: PersistentVolume
apiVersion: v1
metadata:
 name: task-pv-volume
 labels:
 type: local
spec:
 storageClassName: manual
 capacity:
 storage: 10Gi
 accessModes:
 - ReadWriteOnce
 hostPath:
 path: "/mnt/data"
```



# **persistentVolumeClaim**

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
 name: task-pv-claim
spec:
 storageClassName: manual
 accessModes:
 - ReadWriteOnce
 resources:
 requests:
 storage: 3Gi
```



# Pod

```
kind: Pod
apiVersion: v1
metadata:
 name: task-pv-pod
spec:
 containers:
 - name: task-pv-container
 image: nginx
 ports:
 - containerPort: 80
 volumes:
 - name: task-pv-storage
 persistentVolumeClaim:
 claimName: task-pv-claim
 volumeMounts:
 - mountPath: "/usr/share/nginx/html"
 name: task-pv-storage
```

# TP : volumes

